

Deep Generative Models

LMS Code 013672598

Dr. Zulqarnain Khan


Slides from:
Stefano Ermon

Introduction

Challenge: understand complex, unstructured inputs



Computer Vision



Speech




Natural Language Processing



Robotics

Introduction



Richard Feynman: “*What I cannot create, I do not understand*”

Generative modeling: “*What I understand, I can **create***”


Generative Modeling: Computer Graphics

How to generate natural images with a computer?

High level
description


```
Cube(color=blue, position=(x,y,z), size=...)  
Cylinder(color=red, position=(x',y',z'), size=..)
```

Generation (graphics)



Inference (vision as
inverse graphics)

Raw sensory
outputs



Many of our models will have **similar structure (generation + inference)**

Statistical Generative Models

Statistical generative models are **learned from data**




...



Data
(e.g., images of bedrooms)

Prior Knowledge
(e.g., physics, materials, ..)


Priors are always necessary, but there is a spectrum



Statistical Generative Models

A statistical generative model is a **probability distribution** $p(x)$

- **Data:** samples (e.g., images of bedrooms)
- **Prior knowledge:** parametric form (e.g., Gaussian?), loss function (e.g., maximum likelihood?), optimization algorithm, etc.



It is generative because **sampling from $p(x)$ generates new images**




...




Building a simulator for the data generating process




Building a simulator for the data generating process




Building a simulator for the data generating process




Building a simulator for the data generating process




Building a simulator for the data generating process



Building a simulator for the data generating process



Data generation in the real world



Generate

Generative model
of realistic images





Stroke paintings to realistic images
[Meng, He, Song, et al., ICLR 2022]

“Ace of Pentacles”




Generate

Generative model
of paintings




Language-guided artwork creation
<https://chainbreakers.kath.io> @RiversHaveWings

Solving inverse problems with generative models



Generative model
of medical images

Generate


A yellow arrow pointing from the generative model diagram towards the reconstructed brain slices.

Medical image reconstruction
[Song et al., ICLR 2022]

Outlier detection with generative models



High
probability
→



Low
probability
←



Generative model
of traffic signs



Outlier detection
[Song et al., ICLR 2018]

Progress in Generative Models of Images -- GANs




2014



2015



2016



2017



2018

Ian Goodfellow, 2019

Progress in Generative Models of Images – Diffusion Models



Text2Image Diffusion Models

User input:

An astronaut riding a horse



Text2Image Diffusion Models

User input:

A perfect Pakistani meal



Text2Image Diffusion Models

User input:

泰迪熊穿着戏服，站在太和殿前唱京剧

A teddy bear, wearing a costume, is standing in front of the Hall of Supreme Harmony and singing Beijing opera



A model of a cafe adorned with indoor plants. Wooden beams crisscross above, and a cold brew station stands out with tiny bottles and glasses




Progress in Inverse Problems

$P(\text{high resolution} \mid \text{low resolution})$




Menon et al, 2020

$P(\text{full image} \mid \text{mask})$



Liu al, 2018


$P(\text{color image} \mid \text{greyscale})$



Antic, 2020


Progress in Inverse Problems

User input:





Progress in Inverse Problems

Stroke Painting to Image




Stroke-based Editing



Input


Output

Progress in Inverse Problems



Kawar et al., 2023

Medical image reconstruction



Cross-sectional image



\mathbf{x} | \mathbf{y}


Sparse-view
computed
tomography
(CT)





Forward model $p(\mathbf{y} | \mathbf{x})$ is given by physical simulation

WaveNet

Generative model of speech signals



Text to Speech



van den Oord et al, 2016c



Diffusion Text2Speech

Generative model of speech signals



Audio Super Resolution

Conditional generative model $P(\text{high-res signal} \mid \text{low-res audio signal})$



Kuleshov et al., 2017

Language Generation

Custom prompt

To get an A+ in deep generative models, students have to

Completion

To get an A+ in deep generative models, students have to be willing to work with problems that are a whole lot more interesting than, say, the ones that most students work on in class. If you're a great student, the question above can be avoided and you'll be able to do great work, but if you're not, you will need to go beyond the basics before getting good.

Now to be clear, this advice is not just for the deep-learning crowd; it is good advice for any student who is taking his or her first course in machine learning.

The key point is that if you have a deep, deep brain of a computer scientist, that's just as important to you.

$P(\text{next word} \mid \text{previous words})$

Radford et al., 2019
Demo from talktotransformer.com

Language Generation – ChatGPT

You can try the same prompt in ChatGPT, notice the improvement!

Machine Translation

Conditional generative model $P(\text{ English text} | \text{ Chinese text})$





Figure from Google AI research blog.

Code Generation



```
parse_expenses.py  -go write_sql.go  ts sentiment.ts  addresses.rb

1 import datetime
2
3 def parse_expenses(expenses_string):
4     """Parse the list of expenses and return the list of triples (date, va|
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
```

Video Generation

Suddenly, the walls of the embankment broke and there was a huge flood



Video Generation

a couple sledding down a snowy hill on a tire
roman chariot style



Video Generation




Imitation Learning

Conditional generative model $P(\text{actions} \mid \text{past observations})$




Li et al., 2017




Janner et al., 2022

Molecule generation




DeepFakes



Roadmap and Key Challenges

- **Representation:** how do we model the joint distribution of many random variables?
 - Need compact representation
- **Learning:** what is the right way to compare probability distributions?



- **Inference:** how do we invert the generation process (e.g., vision as inverse graphics)?
 - Unsupervised learning: recover high-level descriptions (features) from raw data

Syllabus

- Fully observed likelihood-based models
 - Autoregressive
 - Flow-based models
- Latent variable models
 - Variational learning
 - Inference amortization
 - Variational autoencoder
- Implicit generative models
 - Two sample tests, embeddings, F-divergences
 - Generative Adversarial Networks
- Energy Based Models
- Score-based Diffusion Generative Models
- Learn about algorithms, theory & applications

Prerequisites

- Basic knowledge about machine learning.
- Basic knowledge of probabilities and calculus:
 - Gradients, gradient-descent optimization, backpropagation
 - Random variables, independence, conditional independence
 - Bayes rule, chain rule, change of variables formulas
- Proficiency in some programming language, preferably Python, required.

Logistics

- There is no required textbook. Reading materials and course notes will be provided.
- Suggested Reading: *Deep Learning* by Ian Goodfellow, Yoshua Bengio, Aaron Courville. Online version available free [here](#).

