

# Deep Generative Models (Background)

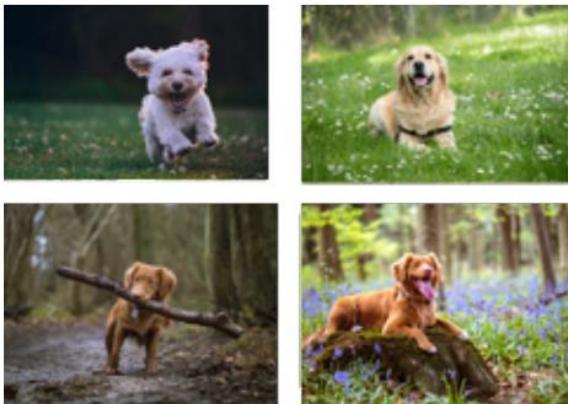
Dr. Zulqarnain Khan

Slides from : Stefano Erman

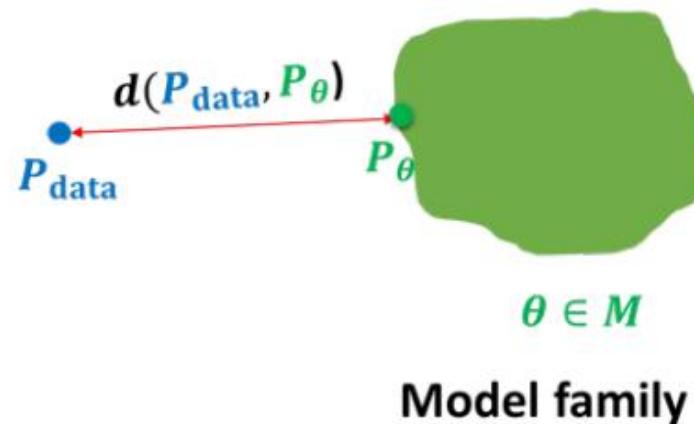
# Overview

- What is a generative model
- Representing probability distributions
  - Curse of dimensionality (what makes generative modeling hard?)
  - Crash course on graphical models (Bayesian networks) – to get familiar with how to represent joint in the form conditionals
  - Generative vs discriminative models – why do generative
  - Neural models – from generative to *deep* generative

# Learning a Generative Model



$\mathbf{x}_i \sim P_{\text{data}}$   
 $i = 1, 2, \dots, n$

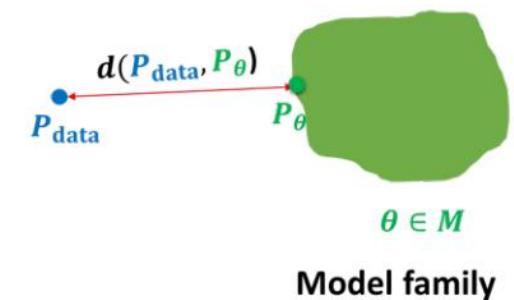


# Learning a Generative Model

- We are given a training set of examples, e.g., images of dogs
- We want to learn a probability distribution  $p(x)$  over images ‘x’ such that we can do:
  - **Generation:** If we sample  $x_{new} \sim p(x)$ ,  $x_{new}$  should look like a dog (sampling)
  - **Density Estimation:**  $p(x)$  should be high if ‘x’ looks like a dog, and low otherwise (anomaly detection)
  - **Unsupervised representation learning:** We should be able to learn what these images have in common, e.g., ears, tail, etc. (features) – possible with at least some of the generative models.
- First question: how to represent  $p(x)$



$$x_i \sim P_{\text{data}} \\ i = 1, 2, \dots, n$$



# Basic Discrete Distributions

- Bernoulli distribution: (biased) coin flip
  - $D = \{\text{Heads}, \text{Tails}\}$
  - Specify  $P(X = \text{Heads}) = p$ . Then  $P(X = \text{Tails}) = 1 - p$ .
  - Write:  $X \sim \text{Ber}(p)$
  - Sampling: flip a (biased) coin
- Categorical distribution: (biased) m-sided dice
  - $D = \{1, \dots, m\}$
  - Specify  $P(Y = i) = p_i$ , such that  $\sum p_i = 1$
  - Write:  $Y \sim \text{Cat}(p_1, \dots, p_m)$
  - Sampling: roll a (biased) die

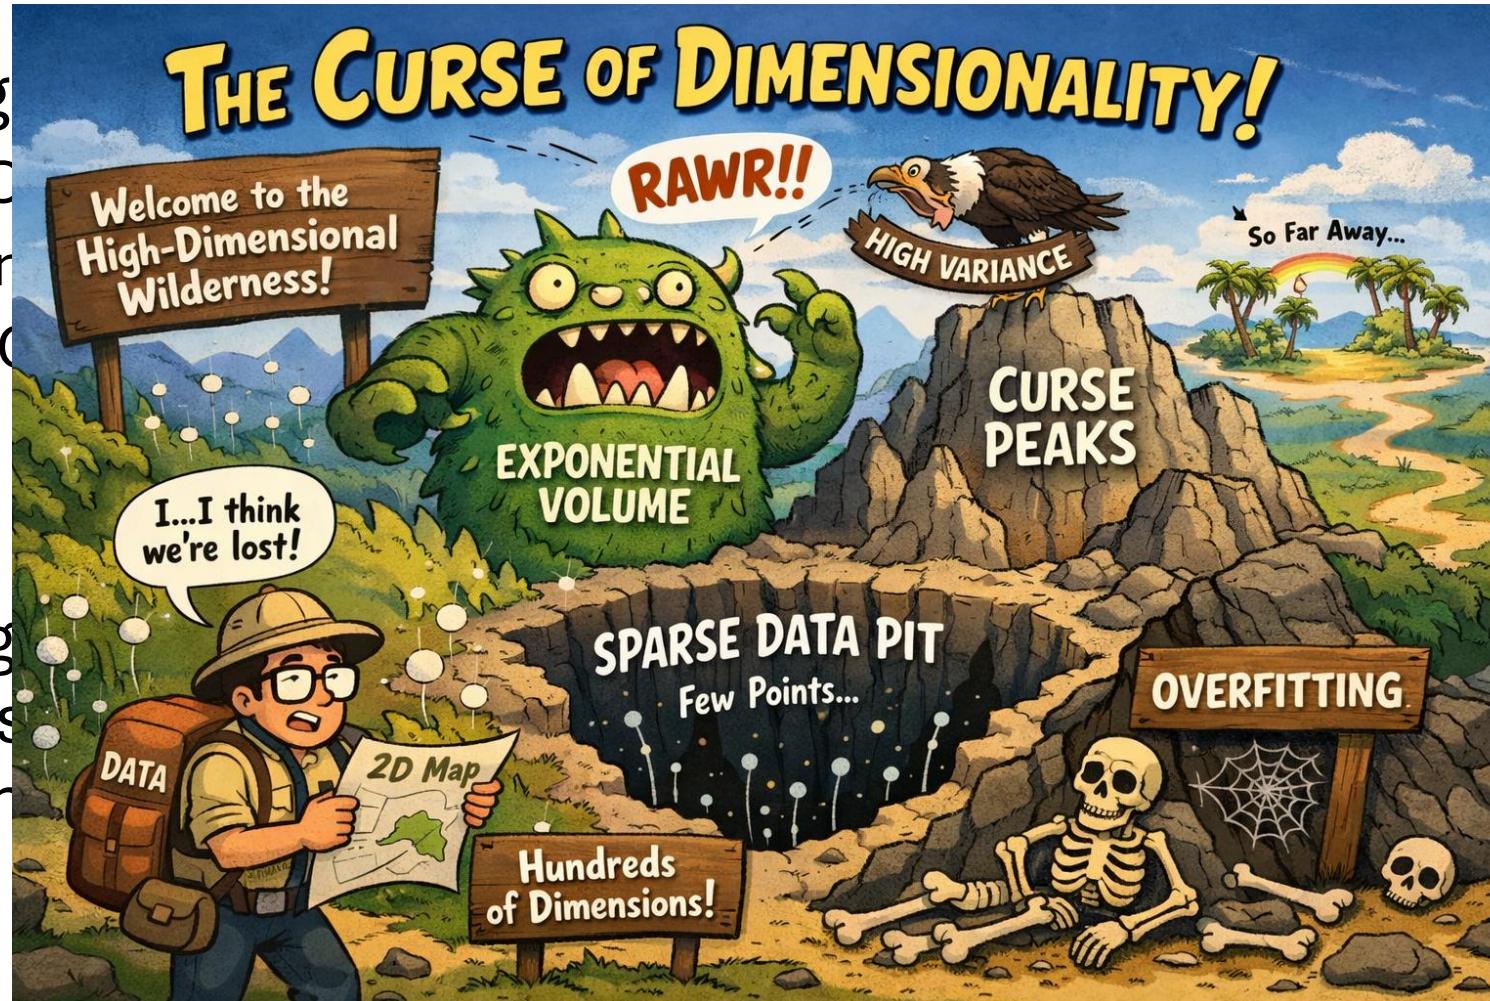
# Example of joint distribution

Modeling

- Red C
- Green
- Blue C

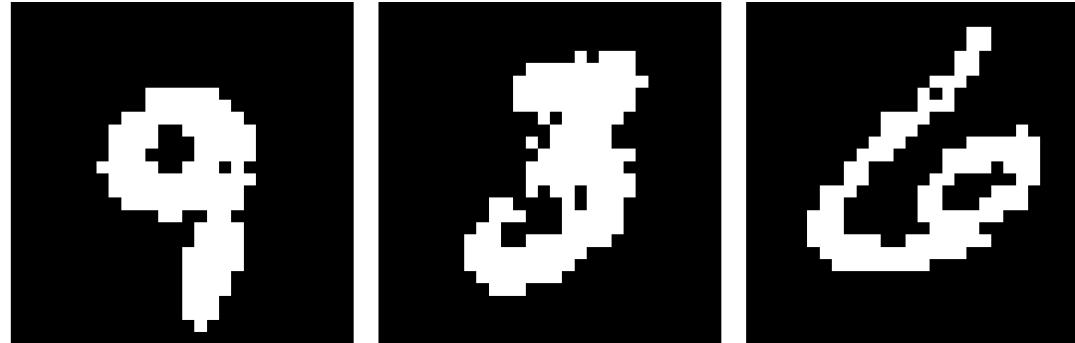
Sampling  
generates  
specify the

variables:



randomly  
we need to

# Example of joint distribution

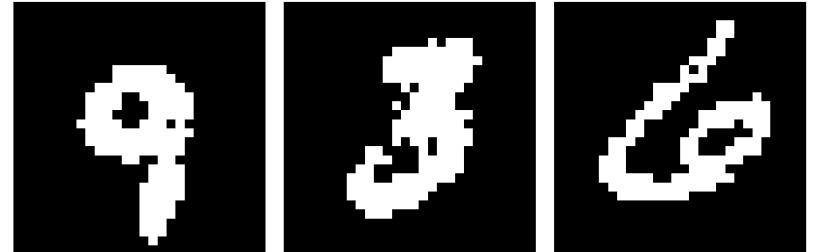


- Suppose  $X_1, \dots, X_n$  (n pixels) are binary (Bernoulli) random variables i.e.  
 $Val(X_i) = \{0, 1\} = \{Black, White\}$
- How many possible images?
- Sampling from  $p(x_1, \dots, x_n)$  generates an image
- How many parameters to specify the joint distribution  $p(x_1, \dots, x_n)$  over n binary pixels?
- These are really big numbers even with black and white pixels!
- We clearly can't do this in full generality. **Solution: Make assumptions!**

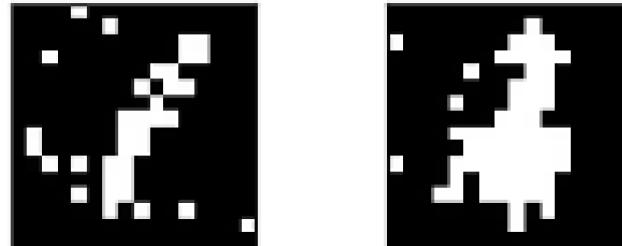
# Structure through independence

- If  $X_1, \dots, X_n$  are independent, then

$$p(x_1, x_2, \dots, x_n) = ?$$



- How many possible images?
- How many parameters now?
- However, independence assumption is too strong. Model not likely to be useful. We are choosing each pixel independently when we sample from this model, and will get:



# Two Important Rules

**Chain rule** Let  $S_1, \dots, S_n$  be events,  $p(S_i) > 0$ .

$$p(S_1 \cap S_2 \cap \dots \cap S_n) = p(S_1)p(S_2 | S_1) \cdots p(S_n | S_1 \cap \dots \cap S_{n-1})$$

**Bayes' rule** Let  $S_1, S_2$  be events,  $p(S_1) > 0$  and  $p(S_2) > 0$ .

$$p(S_1 | S_2) = \frac{p(S_1 \cap S_2)}{p(S_2)} = \frac{p(S_2 | S_1)p(S_1)}{p(S_2)}$$

# Structure through conditional independence

- Using Chain Rule

$$p(x_1, \dots, x_n) = p(x_1)p(x_2 | x_1)p(x_3 | x_1, x_2) \cdots p(x_n | x_1, \dots, x_{n-1})$$

- How many parameters? Still  $2^n - 1$  (why?)

- $p(x_1)$  requires 1 parameter
- But,  $p(x_2 | x_1)$  requires 2 (one each for  $x_1 = 0, 1$ ), and so on...

- Still exponential, chain rule alone doesn't help

- Now suppose  $X_{i+1} \perp X_1, \dots, X_{i-1} | X_i$  (aka Markov), then

$$\begin{aligned} p(x_1, \dots, x_n) &= p(x_1)p(x_2 | x_1)p(x_3 | \cancel{x_1}, x_2) \cdots p(x_n | \cancel{x_1}, \dots, \cancel{x_{n-1}}) \\ &= p(x_1)p(x_2 | x_1)p(x_3 | x_2) \cdots p(x_n | x_{n-1}) \end{aligned}$$

- This gets us to  $2n - 1$  parameters, exponential reduction! You are essentially saying that you only care about the previous variable (e.g. previous word).

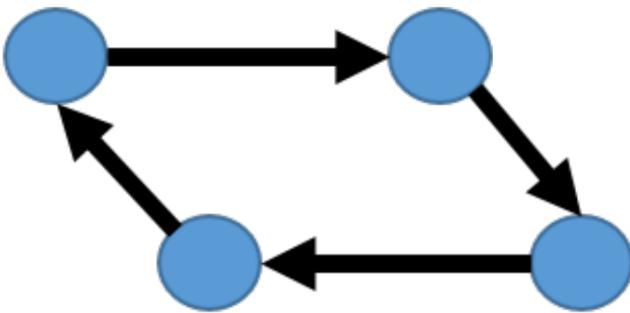
# Bayes Network: General Idea

- Use conditional parameterization (instead of joint parameterization)
- For each random variable  $X_i$ , specify  $p(x_i|\mathbf{x}_{A_i})$  for set  $\mathbf{x}_{A_i}$  of random variables
- Then get joint parametrization as

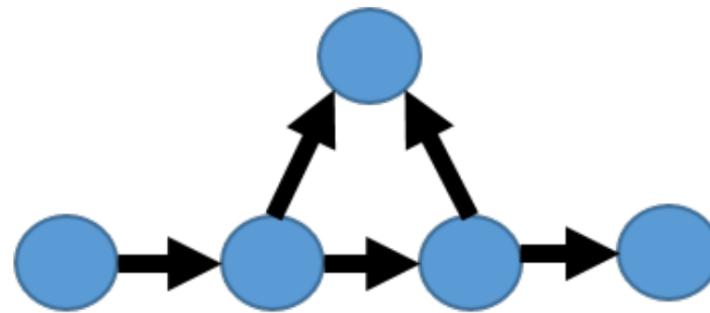
$$p(x_1, \dots, x_n) = \prod_i p(x_i|\mathbf{x}_{A_i})$$

- Has to be a valid probability (correspond graph needs to be a DAG)

# Bayes Network: General Idea



Directed cycle



DAG

DAG stands for Directed Acyclic Graph

# Bayesian networks

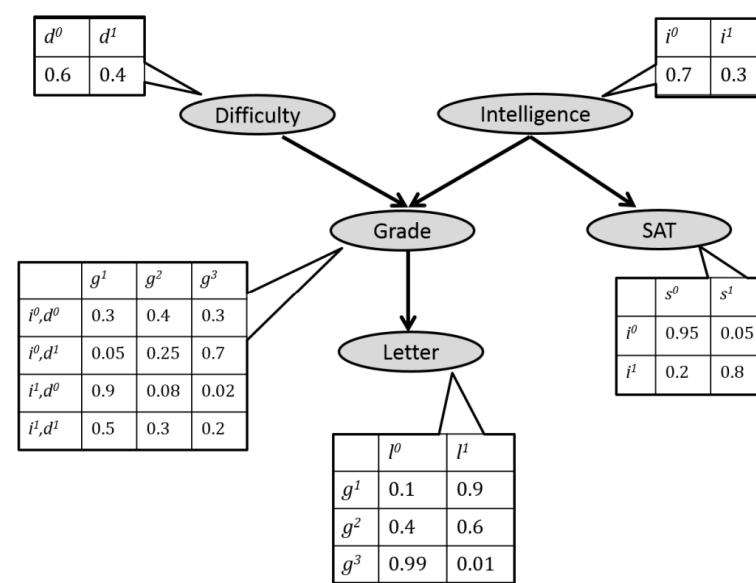
- A Bayesian network is specified by a directed acyclic graph (DAG)  $G = (V, E)$  with:
  - One node  $i \in V$  for each random variable  $X_i$
  - One conditional probability distribution (CPD) per node,  $p(x_i | \mathbf{x}_{Pa(i)})$  specifying the variable's probability conditioned on its parents' values
- Graph  $G = (V, E)$  is called the structure of the Bayesian Network
- Defines a joint distribution:

$$p(x_1, \dots, x_n) = \prod_{i \in V} p(x_i | \mathbf{x}_{Pa(i)})$$

- Economical representation: exponential in  $|Pa(i)|$ , not  $|V|$

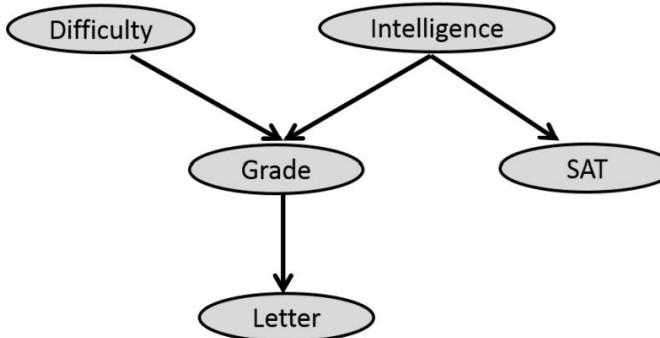
# Example

- Consider the following Bayesian network:



- What is its joint distribution? How many parameters?

# Bayesian network structure implies conditional independencies!



- The joint distribution corresponding to the above BN factors as

$$p(d, i, g, s, l) = p(d)p(i)p(g | i, d)p(s | i)p(l | g)$$

- Generally, by the chain rule, any distribution can be written as

$$p(d, i, g, s, l) = p(d)p(i | d)p(g | i, d)p(s | i, d, g)p(l | g, d, i, s)$$

- Here, we are assuming the following additional independencies:

$$D \perp I, \quad S \perp \{D, G\} | I, \quad L \perp \{I, D, S\} | G$$

# Summary

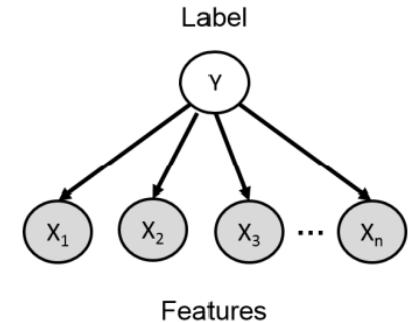
- Bayesian networks given by  $(G, P)$  where  $P$  is specified as a set of local conditional probability distributions associated with graph  $G$ 's nodes
- Efficient representation using a graph-based data structure
- Computing the probability of any assignment is obtained by multiplying CPDs
- Can sample from the joint by sampling from the CPDs according to the DAG ordering
- Can identify some conditional independence properties by looking at graph properties
- In this class, graphical models will be simple (e.g., only 2 or 3 random vectors)
- **Next: generative vs. discriminative**

# Naive Bayes for single label prediction

- Classify e-mails as spam ( $Y = 1$ ) or not spam ( $Y = 0$ )
  - Let  $1 : n$  index the words in our vocabulary (e.g., English)
  - $X_i = 1$  if word  $i$  appears in an e-mail, and 0 otherwise
  - E-mails are drawn according to some distribution  $p(Y, X_1, \dots, X_n)$
- Suppose that the words are conditionally independent given  $Y$ .

Then,

$$p(y, x_1, \dots, x_n) = p(y) \prod_{i=1}^n p(x_i | y)$$



Estimate parameters from training data. Predict with Bayes rule:

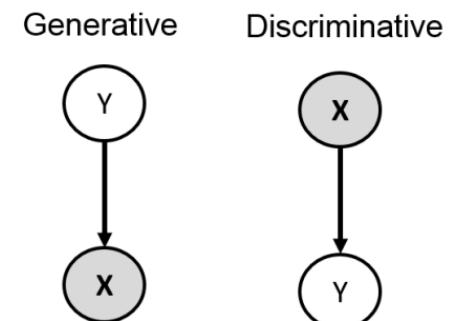
$$p(Y = 1 | x_1, \dots, x_n) = \frac{p(Y = 1) \prod_{i=1}^n p(x_i | Y = 1)}{\sum_{y=\{0,1\}} p(Y = y) \prod_{i=1}^n p(x_i | Y = y)}$$

# Naive Bayes for single label prediction

- Did we make reasonable assumptions in the previous slide?
- Answer: depends. All models are “wrong”, but many are nonetheless useful

# Discriminative versus generative models

- Using chain rule  $p(Y, X) = p(X | Y)p(Y) = p(Y | X)p(X)$ . Corresponding Bayesian networks:



- However, **suppose if all we need for prediction is  $p(Y | X)$**
- In the left model, we need to specify/learn both  $p(Y)$  and  $p(X | Y)$ , then compute  $p(Y | X)$  via Bayes rule
- In the right model, it suffices to estimate just the conditional distribution  $p(Y | X)$ 
  - We never need to model/learn/use  $p(X)$ !
  - Called a discriminative model because it is only useful for discriminating  $Y$ 's label when given  $X$

# Discriminative versus generative models

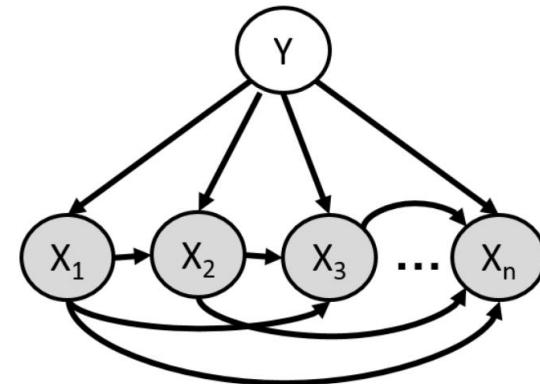
- If all we care about is being able to predict a label -> discriminative
- If we also care about the relationships of all the variables x and y in the data -> generative

# Discriminative versus generative models

Since  $\mathbf{X}$  is a random vector, chain rule will give:

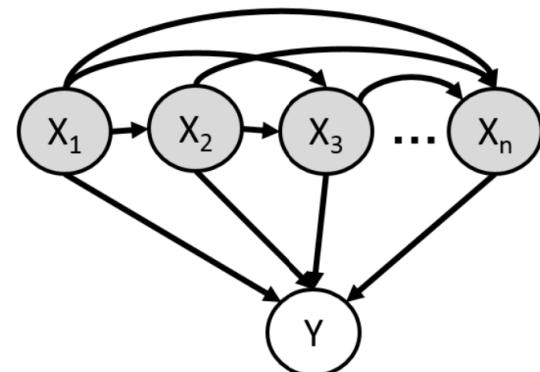
$$p(Y, \mathbf{X}) = p(Y)p(X_1 | Y)p(X_2 | Y, X_1) \cdots p(X_n | Y, X_1, \dots, X_{n-1})$$

Generative



Discriminative

$$p(Y, \mathbf{X}) = p(X_1)p(X_2 | X_1)p(X_3 | X_1, X_2) \cdots p(Y | X_1, \dots, X_{n-1}, X_n)$$



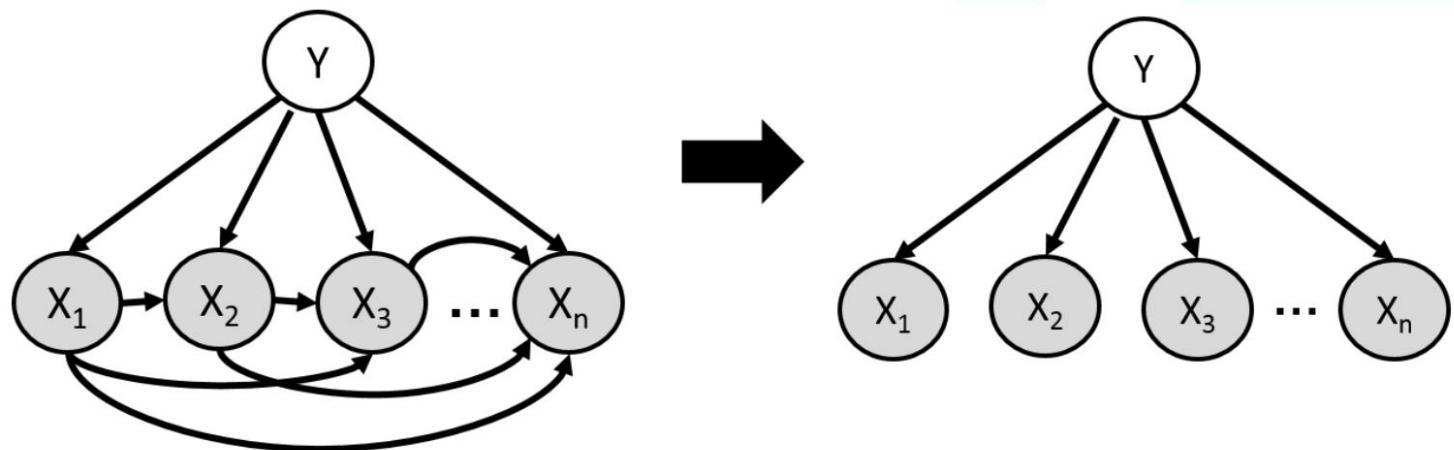
# Discriminative versus generative models

We must make the following choices:

- In the generative model,  $p(Y)$  is simple, but how do we parameterize  $p(X_i | X_{pa(i)}, Y)$ ?
- In the discriminative model, how do we parameterize  $p(Y | X)$ ?  
Here we assume we don't care about modeling  $p(X)$  because  $X$  is always given to us in a classification problem

# Naive Bayes

- For the generative model, assume that  $X_i \perp X_{-i} | Y$  (naive Bayes)



# Logistic regression

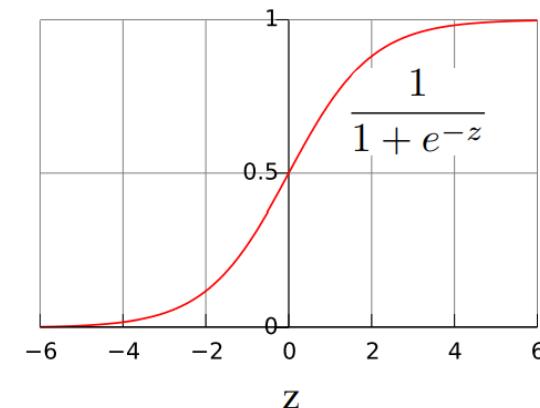
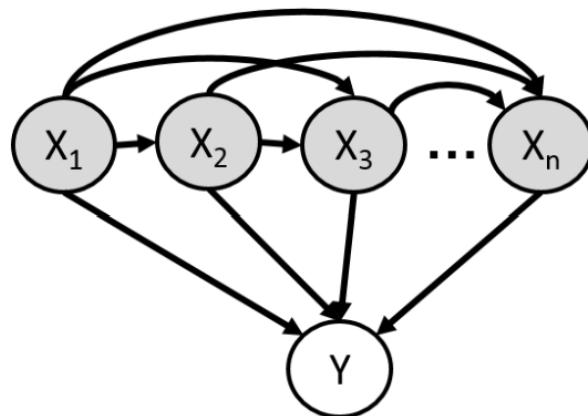
- For the discriminative model, assume that

$$p(Y = 1 \mid \mathbf{x}; \boldsymbol{\alpha}) = f(\mathbf{x}, \boldsymbol{\alpha})$$

- Not represented as a table anymore. It is a parameterized function of  $\mathbf{x}$  (regression)
  - Has to be between 0 and 1
  - Depend in some simple but reasonable way on  $x_1, \dots, x_n$
  - Completely specified by a vector  $\boldsymbol{\alpha}$  of  $n + 1$  parameters (compact representation)

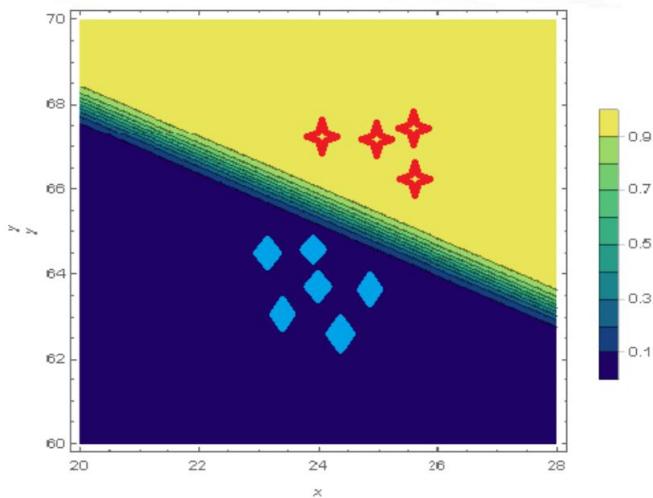
# Logistic Regression

- **Linear dependence:** Let  $z(\alpha, x) = \alpha_0 + \sum_{i=1}^n \alpha_i x_i$ . Then,  $p(Y = 1 | x; \alpha) = \sigma(z(\alpha, x))$ , where  $\sigma(z) = \frac{1}{1 + e^{-z}}$  is the logistic function.

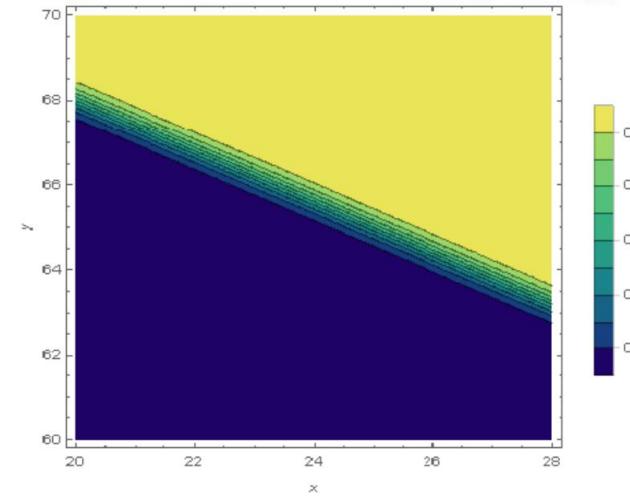


# Logistic Regression

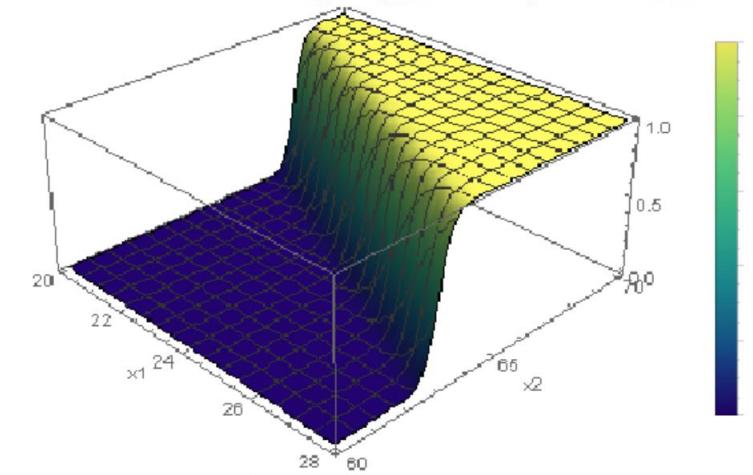
Contours of equal probability defined by  $\alpha$



Probability map defined by  $\alpha$



Probability map defined by  $\alpha$



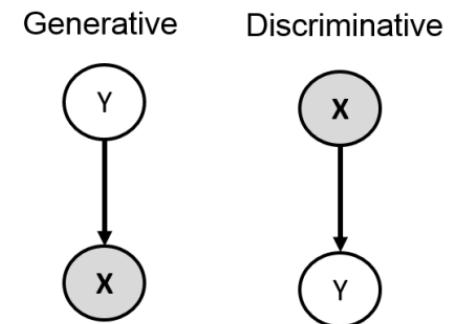
- Decision boundary  $p(Y = 1 | x; \alpha) > 0.5$  is linear in  $x$
- Equal probability contours are straight lines
- Probability rate of change has very specific form (third plot)

# Discriminative models can be powerful

- Make fewer assumptions. e.g. Logistic model does not assume  $X_i \perp X_{-i} | Y$ , unlike naive Bayes.
- This is helpful when there is plenty of data available and you are only interested in prediction.
- This can make a big difference in many applications.

# Generative models are still very useful

- Using chain rule  $p(Y, X) = p(X | Y)p(Y) = p(Y | X)p(X)$ .  
Corresponding Bayesian networks:



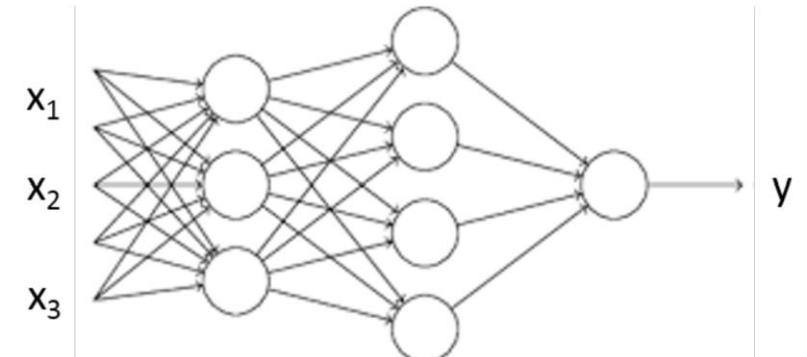
- Using a discriminative conditional model is only possible when  $X$  is always observed.
  - When some  $X_i$  variables are unobserved, the generative model allows us to compute  $p(Y|X_{evidence})$  by marginalizing over the unseen variables
  - Priors can help when there isn't enough data.

# Neural Models

- In discriminative models, we assume that

$$p(Y = 1 \mid \mathbf{x}; \boldsymbol{\alpha}) = f(\mathbf{x}, \boldsymbol{\alpha})$$

- **Linear dependence:** Let  $z(\boldsymbol{\alpha}, \mathbf{x}) = \alpha_0 + \sum_{i=1}^n \alpha_i x_i$ . Then,  $p(Y = 1 \mid \mathbf{x}; \boldsymbol{\alpha}) = \sigma(z(\boldsymbol{\alpha}, \mathbf{x}))$ , where  $\sigma(z) = 1/(1 + e^{-z})$  is the logistic function.  
**This dependence might be too simple.**
- **Non-Linear dependence:** Let  $h(A, b, \mathbf{x}) = f(A\mathbf{x} + b)$  be a non-linear transformation of the inputs. Which makes  $p(Y = 1 \mid \mathbf{x}; \boldsymbol{\alpha}, \mathbf{A}, \mathbf{b}) = f(\alpha_0 + \sum_{i=1}^h \alpha_i h_i)$ 
  - More flexible
  - More parameters:  $\mathbf{A}$ ,  $\mathbf{b}$ ,  $\boldsymbol{\alpha}$
  - Can repeat multiple times to get a neural network



# Bayesian networks vs neural models

- Using Chain Rule (fully general)

$$p(x_1, x_2, x_3, x_4) = p(x_1)p(x_2 | x_1)p(x_3 | x_1, x_2)p(x_4 | x_1, x_2, x_3)$$

- Bayes Net (assumes conditional independencies)

$$p(x_1, x_2, x_3, x_4) \approx p(x_1)p(x_2 | x_1)p(x_3 | \cancel{x_1}, x_2)p(x_4 | \cancel{x_1}, \cancel{x_2}, \cancel{x_3})$$

- Neural Models (assume specific functional form for the conditionals. A sufficiently deep neural net can approximate any function.)

$$p(x_1, x_2, x_3, x_4) \approx p(x_1)p(x_2 | x_1)p_{\text{Neural}}(x_3 | x_1, x_2)p_{\text{Neural}}(x_4 | x_1, x_2, x_3)$$

# Continuous variables

- If  $X$  is a continuous random variable (scalar), we can usually represent it using its **probability density function**. Typically consider parameterized densities, like Gaussian, or Uniform.
- Correspondingly for a continuous random vector  $X$ , we can usually represent it using its ***joint* probability density function**. E.g.

$$\text{Gaussian: if } p_X(x) = \frac{1}{\sqrt{(2\pi)^n |\Sigma|}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1} (x - \mu)\right)$$

- **Chain rule, Bayes rule, etc all still apply.** For example,

$$p_{X,Y,Z}(x, y, z) = p_X(x)p_{Y|X}(y | x)p_{Z|{\{X,Y\}}}(z | x, y)$$

# Continuous variables

- This means we can still use Bayesian networks with continuous (and discrete) variables. Examples:
- **Mixture of 2 Gaussians**: Bayes net  $Z \rightarrow X$  with factorization  $p_{Z,X}(z,x) = p_Z(z)p_{X|Z}(x|z)$  and,
  - $Z \sim \text{Bernoulli}(p)$ , i.e.  $z$  is binary
  - $X|Z=0 \sim N(\mu_0, \sigma_0)$ ,  $X|Z=1 \sim N(\mu_1, \sigma_1)$
  - The parameters are  $p, \mu_0, \sigma_0, \mu_1, \sigma_1$
- Bayes net  $Z \rightarrow X$  with factorization  $p(z,x) = p(z)p(x|z)$ 
  - $Z \sim U(a, b)$
  - $X | (Z=z) \sim N(z, \sigma)$
  - The parameters are  $a, b, \sigma$