

Maximum Likelihood Learning

Dr. Zulqarnain Khan


(Slides from: Stefano Ermon)

Learning a Generative Model

- We are given a training set of examples, e.g., images of dogs
- We want to learn a probability distribution $p(x)$ over images ‘x’ such that we can do:
 - **Generation:** If we sample $x_{new} \sim p(x)$, x_{new} should look like a dog (sampling)
 - **Density Estimation:** $p(x)$ should be high if ‘x’ looks like a dog, and low otherwise (anomaly detection)
 - **Unsupervised representation learning:** We should be able to learn what these images have in common, e.g., ears, tail, etc. (features) – possible with at least some of the generative models.
- First question: how to represent $p(x)$ – today we’ll talk about AR models
- Second question: how to learn it! ←



$$x_i \sim P_{\text{data}} \\ i = 1, 2, \dots, n$$



Model family

Setting

- Let's assume that the domain is governed by some underlying distribution P_{data}
- We are given a dataset D of m samples from P_{data}
 - Each sample is an assignment of values to (a subset of) the variables, e.g., $(X_{bank} = 1, X_{dollar} = 0, \dots, Y = 1)$ or pixel intensities
- The standard assumption is that the data instances are **independent and identically distributed (IID)**
- We are also given a **family of models M**, and our task is to learn some “good” distribution in this set:
 - For example, M could be all Bayes nets with a given graph structure, for all possible choices of the CPD tables
 - For example, a FVSBN for all possible choices of the logistic regression parameters , θ = concatenation of all logistic regression coefficients

Goal of learning

- The goal of learning is to return a model P_θ that precisely captures the distribution P_{data} from which our data was sampled
- This is in general not achievable because of
 - limited data only provides a rough approximation of the true underlying distribution
 - computational reasons
- Example. Images represented with a vector X of 784 binary variables (black vs. white pixel). 2^{784} possible images/states. Even 10million examples is extremely sparse coverage.
- We want to select P_θ to construct the “best” approximation to the underlying distribution P_{data}
- What is “best”?

What is “best”?


- This depends on what we want to do
-
- **Density estimation:** we are interested in the full distribution (so later we can compute whatever conditional probabilities we want)
 - **Specific prediction tasks:** we are using the distribution to make a prediction
 - Is this email spam or not?
 - Structured prediction: Predict next frame in a video, or caption given an image 3
 - **Structure or knowledge discovery:** we are interested in the model itself (e.g. causal models)
 - How do some genes interact with each other?
 - What causes cancer?

Learning as density estimation

- We want to learn the full distribution so that later we can answer any probabilistic inference query
- In this setting we can view the learning problem as density estimation
- We want to construct P_θ as “close” as possible to P_{data} (recall we assume we are given a dataset D of samples from P_{data})
- How do we evaluate “closeness”?



$$\mathbf{x}_i \sim P_{data} \\ i = 1, 2, \dots, n$$



KL-divergence (review)

- How should we measure distance between distributions?
- The Kullback-Leibler divergence (KL-divergence) between two distributions p and q is defined as

$$D(p||q) = \sum_x p(x) \log \frac{p(x)}{q(x)}$$

- $D(p||q) \geq 0, \forall p, q$ with equality if and only if $p = q$. Proof:

$$\mathbf{E}_{\mathbf{x} \sim p} \left[-\log \frac{q(\mathbf{x})}{p(\mathbf{x})} \right] \geq -\log \left(\mathbf{E}_{\mathbf{x} \sim p} \left[\frac{q(\mathbf{x})}{p(\mathbf{x})} \right] \right) = -\log \left(\sum_{\mathbf{x}} p(\mathbf{x}) \frac{q(\mathbf{x})}{p(\mathbf{x})} \right) = 0$$

- Notice that KL-divergence is **asymmetric**, i.e., $D(p \parallel q) \neq D(q \parallel p)$
- Measures the expected number of extra bits required to describe samples from $p(\mathbf{x})$ using a compression code based on q instead of p

KL-divergence (review)

- To compress, it is useful to know the probability distribution the data is sampled from
- For example, let X_1, \dots, X_{100} be samples of an unbiased coin. Roughly 50 heads and 50 tails. Optimal compression scheme is to record heads as 0 and tails as 1. In expectation, use 1 bit per sample, and cannot do better
- Suppose the coin is biased, and $P[H] \gg P[T]$. Then it's more efficient to use fewer bits on average to represent heads and more bits to represent tails, e.g.
 - Use a short sequence of bits to encode HHHH (common) and a long sequence for TTTT (rare).
 - Like Morse code: E = •, A = •–, Q = –– •–
- KL-divergence: if your data comes from p , but you use a scheme optimized for q , the divergence $D_{KL}(p||q)$ is the number of extra bits you'll need on average

Learning as density estimation

- We want to learn the full distribution so that later we can answer any probabilistic inference query
- In this setting we can view the learning problem as density estimation
- We want to construct P_{θ} as "close" as possible to P_{data} (recall we assume we are given a dataset D of samples from P_{data}) How do we evaluate "closeness"? KL-divergence is one possibility:

$$D(P_{data} || P_{\theta}) = \mathbb{E}_{\mathbf{x} \sim P_{data}} \left[\log \left(\frac{P_{data}(\mathbf{x})}{P_{\theta}(\mathbf{x})} \right) \right] = \sum_{\mathbf{x}} P_{data}(\mathbf{x}) \log \frac{P_{data}(\mathbf{x})}{P_{\theta}(\mathbf{x})}$$

- $D(P_{data} || P_{\theta}) = 0$ iff the two distributions are the same.
- It measures "compression loss" (in bits) of using P_{θ} instead of P_{data}

Expected log-likelihood

- We can simplify this somewhat:

$$\begin{aligned} D(P_{\text{data}} || P_{\theta}) &= \mathbb{E}_{x \sim P_{\text{data}}} \left[\log \left(\frac{P_{\text{data}}(x)}{P_{\theta}(x)} \right) \right] \\ &= \mathbb{E}_{x \sim P_{\text{data}}} [\log P_{\text{data}}(x)] - \mathbb{E}_{x \sim P_{\text{data}}} [\log P_{\theta}(x)] \end{aligned}$$

No θ ←

- Then, minimizing KL divergence is equivalent to maximizing the expected log-likelihood

$$\arg \min_{P_{\theta}} D(P_{\text{data}} || P_{\theta}) = \arg \min_{P_{\theta}} -\mathbb{E}_{x \sim P_{\text{data}}} [\log P_{\theta}(x)] = \arg \max_{P_{\theta}} \mathbb{E}_{x \sim P_{\text{data}}} [\log P_{\theta}(x)]$$

- Asks that P_{θ} assign high probability to instances sampled from P_{data} , so as to reflect the true distribution
- Because of log, samples x where $P_{\theta}(x) \approx 0$ weigh heavily in objective
- Although we can now compare models, since we are ignoring $H(P_{\text{data}}) = -\mathbb{E}_{x \sim P_{\text{data}}} [\log P_{\text{data}}(x)]$, we don't know how close we are to the optimum.
- This still involves expectation over P_{data} , we do not know P_{data} !

Maximum likelihood

- Approximate the expected log-likelihood

$$\mathbb{E}_{\mathbf{x} \sim P_{\text{data}}} [\log P_{\theta}(\mathbf{x})]$$

using the empirical log-likelihood:

$$\mathbb{E}_{\mathcal{D}} [\log P_{\theta}(\mathbf{x})] = \frac{1}{|\mathcal{D}|} \sum_{\mathbf{x} \in \mathcal{D}} \log P_{\theta}(\mathbf{x})$$

- Maximum likelihood learning is then:

$$\max_{P_{\theta}} \quad \frac{1}{|\mathcal{D}|} \sum_{\mathbf{x} \in \mathcal{D}} \log P_{\theta}(\mathbf{x})$$

- Equivalently, maximize likelihood of the data

$$P_{\theta}(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}) = \prod_{\mathbf{x} \in \mathcal{D}} P_{\theta}(\mathbf{x})$$

Monte Carlo Estimation (Recap)

1. Express the quantity of interest as the expected value of a random variable.

$$E_{x \sim P}[g(x)] = \sum_x g(x)P(x)$$

2. Generate T samples x^1, \dots, x^T from the distribution P with respect to which the expectation was taken.
3. Estimate the expected value from the samples using:

$$\hat{g}(x^1, \dots, x^T) \triangleq \frac{1}{T} \sum_{t=1}^T g(x^t)$$

where x^1, \dots, x^T are independent samples from P . Note: \hat{g} is a random variable. Why?

Properties of the Monte Carlo Estimate

- Unbiased

$$E_P[\hat{g}] = E_P[g(x)]$$

- Convergence: By law of large numbers

$$\hat{g} = \frac{1}{T} \sum_{t=1}^T g(x^t) \rightarrow E_P[g(x)] \text{ for } T \rightarrow \infty$$

- Variance

$$V_P[\hat{g}] = V_P \left[\frac{1}{T} \sum_{t=1}^T g(x^t) \right] = \frac{V_P[g(x)]}{T}$$

Thus, variance of the estimator can be reduced by increasing the number of samples.


Example

Single variable example: A biased coin

- Two outcomes: heads (H) and tails (T)
- Data set: Tosses of the biased coin, e.g., $D = \{H, H, T, H, T\}$
- Assumption: the process is controlled by a probability distribution $P_{data}(x)$ where $x \in \{H, T\}$
- Class of models M: all probability distributions over $x \in \{H, T\}$.
- Example learning task: How should we choose $P_\theta(x)$ from M if 3 out of 5 tosses are heads in D?

MLE scoring for the coin example

- We represent our model: $P_\theta(x = H) = \theta$, and $P_\theta(x = T) = 1 - \theta$
 - Example $D = \{H, H, T, H, T\}$
 - Likelihood of data = $\prod_i P_\theta(x_i = \theta) \cdot \theta \cdot (1 - \theta) \cdot \theta \cdot (1 - \theta)$



- Optimize for θ which makes D most likely. What is the solution in this case? $\theta = 0.6$, optimization problem can be solved in closed-form

MLE principle for Autoregressive Models

Given an autoregressive model with n variables and factorization

$$P_{\theta}(\mathbf{x}) = \prod_{i=1}^n p_{\text{neural}}(x_i | \mathbf{x}_{<i}; \theta_i)$$

$\theta = (\theta_1, \dots, \theta_n)$ are the parameters of all the conditionals. Training data $D = \{x^{(1)}, \dots, x^{(m)}\}$. Maximum likelihood estimate of the parameters θ ?

- Decomposition of Likelihood function

$$L(\theta, D) = \prod_{j=1}^m P_{\theta}(\mathbf{x}^{(j)}) = \prod_{j=1}^m \prod_{i=1}^n p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i)$$

- Goal: maximize $\arg \max_{\theta} L(\theta, D) = \arg \max_{\theta} \log L(\theta, D)$
- We no longer have a closed form solution

MLE Learning: Gradient Descent

$$L(\theta, \mathcal{D}) = \prod_{j=1}^m P_\theta(\mathbf{x}^{(j)}) = \prod_{j=1}^m \prod_{i=1}^n p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i)$$

Goal: maximize $\arg \max_{\theta} L(\theta, D) = \arg \max_{\theta} \log L(\theta, D)$

$$\ell(\theta) = \log L(\theta, \mathcal{D}) = \sum_{j=1}^m \sum_{i=1}^n \log p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i)$$

1. Initialize $\theta^0 = (\theta_1, \dots, \theta_n)$ at random
2. Compute $\nabla_{\theta} l(\theta)$ (by back propagation)
3. $\theta^{t+1} = \theta^t + \alpha_t \nabla_{\theta} l(\theta)$

Non-convex optimization problem, but often works well in practice

MLE Learning: Stochastic Gradient Descent

$$\ell(\theta) = \log L(\theta, \mathcal{D}) = \sum_{j=1}^m \sum_{i=1}^n \log p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i)$$

1. Initialize θ^0 at random
2. Compute $\nabla_\theta(\theta)$ (by back prop)
3. $\theta^{k+1} = \theta^t + \alpha_t \nabla_\theta(\theta)$

What is the gradient with respect to θ_i ?

$$\nabla_{\theta_i} \ell(\theta) = \sum_{j=1}^m \nabla_{\theta_i} \sum_{i=1}^n \log p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i) = \sum_{j=1}^m \nabla_{\theta_i} \log p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i)$$

Each conditional $p_{\text{neural}}(x_i | x_{<i}; \theta_i)$ can be optimized separately if there is no parameter sharing. In practice, parameters θ_i are shared (e.g., NADE, PixelRNN, PixelCNN, etc.)

MLE Learning: Stochastic Gradient Descent

$$\ell(\theta) = \log L(\theta, \mathcal{D}) = \sum_{j=1}^m \sum_{i=1}^n \log p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i)$$

1. Initialize θ^0 at random
2. Compute $\nabla_\theta(\theta)$ (by back prop)
3. $\theta^{k+1} = \theta^t + \alpha_t \nabla_\theta(\theta)$

$$\nabla_\theta \ell(\theta) = \sum_{j=1}^m \sum_{i=1}^n \nabla_\theta \log p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i)$$

What if $m = |\mathcal{D}|$ is huge?

$$\nabla_\theta \ell(\theta) = m \sum_{j=1}^m \frac{1}{m} \sum_{i=1}^n \nabla_\theta \log p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i) = m E_{x^{(j)} \sim \mathcal{D}} \left[\sum_{i=1}^n \nabla_\theta \log p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i) \right]$$

Monte Carlo: Sample $x^{(j)} \sim \mathcal{D}; \nabla_\theta \ell(\theta) \approx m \sum_{i=1}^n \nabla_\theta \log p_{\text{neural}}(x_i^{(j)} | \mathbf{x}_{<i}^{(j)}; \theta_i)$

Empirical Risk and Overfitting

- Empirical risk minimization can easily **overfit** the data
 - Extreme example: The data is the model (memorize all training data).
- **Generalization**: the data is a sample, usually there is vast amount of samples that you have never seen. Your model should generalize well to these “never-seen” samples.
- Thus, we typically restrict the hypothesis space of distributions that we search over




Bias-Variance trade off

- If the hypothesis space is very limited, it might not be able to represent P_{data} , even with unlimited data
 - This type of limitation is called **bias**, as the learning is limited on how close it can approximate the target distribution
- If we select a highly expressive hypothesis class, we might represent better the data
 - When we have small amount of data, multiple models can fit well, or even better than the true model. Moreover, small perturbations on D will result in very different estimates
 - This limitation is call the **variance**.

Bias-Variance trade off

There is an inherent bias-variance trade off when selecting the hypothesis class. Error in learning due to both things: bias and variance.

- Hypothesis space: linear relationship
 - Does it fit well? Underfits
- Hypothesis space: high degree polynomial
 - Overfits
- Hypothesis space: low degree polynomial
 - Right tradeoff




How to avoid overfitting?

- **Hard constraints**, e.g. by selecting a less expressive model family:

- Smaller neural networks
- Weight sharing



$$\mathbf{x}_i \sim P_{\text{data}} \\ i = 1, 2, \dots, n$$



- Soft preference for “simpler” models: **Occam Razor**.
- Augment the objective function with **regularization**:

$$\text{objective}(\mathbf{x}, \mathcal{M}) = \text{loss}(\mathbf{x}, \mathcal{M}) + R(\mathcal{M})$$

- Evaluate **generalization performance** on a held-out validation set

Conditional generative models

- Suppose we want to generate a set of variables Y given some others X, e.g., text to speech
- We concentrate on modeling $p(Y|X)$, and use a conditional loss function
 - $-\log P_\theta(y | x)$.
- Since the loss function only depends on $P_\theta(y | x)$, suffices to estimate the conditional distribution, not the joint



Input: image



Brown horse in
grass field

Output: caption

Recap

- For autoregressive models, it is easy to compute $p_{\theta}(x)$
 - Ideally, evaluate in parallel each conditional $\log p_{neural}(x_i^{(j)} | x_{<i}^{(j)}; \theta_i)$. Not like RNNs.
- Natural to train them via maximum likelihood
- Higher log-likelihood doesn't necessarily mean better looking samples
- Other ways of measuring similarity are possible (Generative Adversarial Networks, GANs)